

```
PS C:\Users\bigmo> wsl -d Ubuntu -- /home/bigmo/sglang-challenge/test_request.sh
wsl: 检测到 localhost 代理配置, 但未镜像到 WSL。NAT 模式下的 WSL 不支持 localhost 代理。
=== GET /v1/models ===
{"object": "list", "data": [{"id": "Qwen/Qwen3-0.6B", "object": "model", "created": 1790395798, "owned_by": "sglang", "root": "Qwen/Qwen3-0.6B", "parent": null, "max_model_len": 40960}]}
=== POST /v1/chat/completions ===
{"id": "f8ca4a459ced4f1ba52b64346efa1a32", "object": "chat.completion", "created": 1790395800, "model": "Qwen/Qwen3-0.6B", "choices": [{"index": 0, "message": {"role": "assistant", "content": "<think>\n好的, 用户让我用一句中文介绍SGLang。首先, 我需要确认SGLang是什么。SGLang可能是指某个特定的编程语言或框架, 但不确定。可能用户指的是SGL (Simplified General Language) 或者某个特定的项目名称。不过, 考虑到中文的常见", "reasoning_content": null, "tool_calls": null}, "logprobs": null, "finish_reason": "length", "matched_stop": null}], "usage": {"prompt_tokens": 17, "total_tokens": 81, "completion_tokens": 64, "prompt_tokens_details": null, "reasoning_tokens": 0}, "metadata": {"weight_version": "default"}}
PS C:\Users\bigmo> |
```

```
(.venv) bigmo@huang:~/sglang-challenge/workload$ python run_mooncake_workload.py
[03] input= 931 output= 23 status=success      TTFT=0.1265s latency=0.6775s
[01] input=1001 output=101 status=success      TTFT=0.1982s latency=2.9218s
[02] input= 894 output= 39 status=success      TTFT=0.1207s latency=1.1099s
[05] input= 893 output= 30 status=success      TTFT=0.0696s latency=0.7472s
[04] input= 912 output= 91 status=success      TTFT=0.0862s latency=2.2733s
[06] input= 894 output= 89 status=success      TTFT=0.0595s latency=2.2723s
[07] input= 901 output=  8 status=success      TTFT=0.0708s latency=0.2445s
[08] input= 932 output= 10 status=success      TTFT=0.072s latency=0.3711s
[09] input= 910 output= 59 status=success      TTFT=0.0928s latency=1.6228s
[10] input= 958 output= 24 status=success      TTFT=0.0609s latency=0.5705s
```

Completed: 10/10 successful

CSV saved to: /home/bigmo/sglang-challenge/workload/workload_results.csv

```
(.venv) bigmo@huang:~/sglang-challenge/workload$ python run_mooncake_workload.py
```

```
[03] input= 931 output= 23 status=success      TTFT=0.1164s latency=0.9031s
[02] input= 894 output= 39 status=success      TTFT=0.1296s latency=1.4928s
[05] input= 893 output= 30 status=success      TTFT=0.0649s latency=0.9639s
[01] input=1001 output=101 status=success      TTFT=0.9854s latency=4.6573s
[04] input= 912 output= 91 status=success      TTFT=0.1043s latency=2.9166s
[07] input= 901 output=  8 status=success      TTFT=0.0881s latency=0.3535s
[06] input= 894 output= 89 status=success      TTFT=0.0741s latency=2.8755s
[08] input= 932 output= 10 status=success      TTFT=0.0753s latency=0.3933s
[10] input= 958 output= 24 status=success      TTFT=0.0887s latency=0.7835s
[09] input= 910 output= 59 status=success      TTFT=0.0855s latency=1.946s
```